

```
In [ ]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
%matplotlib inline
```

```
In [ ]: df=pd.read_csv('/content/drive/MyDrive/project ai/UCI Heart Disease  
Data.csv')  
df
```

	id	age	sex	dataset	cp	trestbps	chol
0	1	63	Male	Cleveland	typical angina	145.0	233.0
1	2	67	Male	Cleveland	asymptomatic	160.0	286.0
2	3	67	Male	Cleveland	asymptomatic	120.0	229.0
3	4	37	Male	Cleveland	non-anginal	130.0	250.0
4	5	41	Female	Cleveland	atypical angina	130.0	204.0
...
915	916	54	Female	VA Long Beach	asymptomatic	127.0	333.0
916	917	62	Male	VA Long Beach	typical angina	NaN	139.0
917	918	55	Male	VA Long Beach	asymptomatic	122.0	223.0
918	919	58	Male	VA Long Beach	asymptomatic	NaN	385.0
919	920	62	Male	VA Long Beach	atypical angina	120.0	254.0

920 rows × 16 columns

```
In [ ]: df.head()
```

	id	age	sex	dataset	cp	trestbps	chol	fbs
0	1	63	Male	Cleveland	typical angina	145.0	233.0	True
1	2	67	Male	Cleveland	asymptomatic	160.0	286.0	False
2	3	67	Male	Cleveland	asymptomatic	120.0	229.0	False
3	4	37	Male	Cleveland	non-anginal	130.0	250.0	False
4	5	41	Female	Cleveland	atypical angina	130.0	204.0	False

```
In [ ]: df.columns.values
```

```
array(['id', 'age', 'sex', 'dataset', 'cp', 'trestbps', 'chol', 'fbs',
       'restecg', 'thalch', 'exang', 'oldpeak', 'slope', 'ca', 'thal',
       'num'], dtype=object)
```

```
In [ ]: df.isna().sum()
```


	0
id	0
age	0
sex	0
dataset	0
cp	0
trestbps	59
chol	30
fbs	90
restecg	2
thalch	55
exang	55
oldpeak	62
slope	309
ca	611
thal	486
num	0

dtype: int64

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 920 entries, 0 to 919
Data columns (total 16 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   id         920 non-null    int64  
 1   age        920 non-null    int64  
 2   sex        920 non-null    object  
 3   dataset    920 non-null    object  
 4   cp         920 non-null    object  
 5   trestbps   861 non-null    float64 
 6   chol       890 non-null    float64 
 7   fbs        830 non-null    object  
 8   restecg    918 non-null    object  
 9   thalch     865 non-null    float64 
 10  exang      865 non-null    object  
 11  oldpeak    858 non-null    float64 
 12  slope      611 non-null    object  
 13  ca          309 non-null    float64 
 14  thal       434 non-null    object  
 15  num         920 non-null    int64  
dtypes: float64(5), int64(3), object(8)
memory usage: 115.1+ KB
```

```
In [ ]: df.hist(bins = 50, grid = False, figsize = (20,15));
```



```
In [ ]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, OneHotEncoder,
StandardScaler, MinMaxScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
```

```
In [ ]: #handle missing value
missing_counts = df.isna().sum()
print(missing_counts[missing_counts > 0])
```

```
      trestbps    59
      chol       30
      fbs        90
      restecg     2
      thalch     55
      exang      55
      oldpeak     62
      slope      309
      ca         611
      thal      486
      dtype: int64
```

```
In [ ]: #Duplicate rows remove
print("Duplicates before:", df.duplicated().sum())
df = df.drop_duplicates().reset_index(drop=True)
print("Shape after dropping duplicates:", df.shape)
```

```
Duplicates before: 0
Shape after dropping duplicates: (920, 16)
```

```
In [ ]: #Correct data type
if 'oldpeak' in df.columns:
    df['oldpeak'] = df['oldpeak'].astype('float64')
    print(df['oldpeak'].value_counts())
```

```
oldpeak
0.0      370
1.0       83
2.0       76
1.5       48
3.0       28
0.5       19
1.2       17
2.5       16
0.8       15
1.4       15
0.6       14
0.2       14
1.6       14
1.8       12
0.4       10
0.1        9
4.0        8
2.6        7
2.8        7
2.2        5
1.3        5
0.3        5
0.7        5
1.9        5
1.1        4
3.6        4
2.4        4
0.9        4
3.4        3
-0.5       2
3.2        2
2.3        2
3.5        2
1.7        2
-1.0       2
4.2        2
2.1        2
3.1        1
2.9        1
6.2        1
5.6        1
3.8        1
-1.1       1
4.4        1
5.0        1
-2.6       1
-0.1       1
```

```
-1.5      1  
-2.0      1  
-0.7      1  
-0.8      1  
-0.9      1  
 3.7      1  
Name: count, dtype: int64
```

```
In [ ]: import pandas as pd
```

```
df = pd.read_csv("/content/drive/MyDrive/project ai/UCI Heart Disease  
Data.csv")  
print("Dataset Loaded Successfully!")
```

Dataset Loaded Successfully!

```
In [ ]: import pandas as pd
from sklearn.preprocessing import LabelEncoder

# Load dataset
df = pd.read_csv("//content/drive/MyDrive/project ai/UCI Heart Disease
Data.csv")

# Categorical columns in your dataset
categorical_cols = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca',
'thal', 'target']

le = LabelEncoder()

for col in categorical_cols:
    if col in df.columns:

        # 1. Fill missing values first
        df[col] = df[col].fillna(df[col].mode()[0])

        # 2. If string → Label Encode it
        if df[col].dtype == 'object':
            df[col] = le.fit_transform(df[col])

        # 3. Now safely convert to int
        df[col] = df[col].astype(int)

print("✓ All categorical columns encoded successfully!")
```

✓ All categorical columns encoded successfully!

```
/tmp/ipython-input-3325805646.py:16: FutureWarning: Downcasting
object dtype arrays on .fillna, .ffill, .bfill is deprecated and
will change in a future version. Call
result.infer_objects(copy=False) instead. To opt-in to the future
behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    df[col] = df[col].fillna(df[col].mode()[0])
/tmp/ipython-input-3325805646.py:16: FutureWarning: Downcasting
object dtype arrays on .fillna, .ffill, .bfill is deprecated and
will change in a future version. Call
result.infer_objects(copy=False) instead. To opt-in to the future
behavior, set `pd.set_option('future.no_silent_downcasting', True)`
    df[col] = df[col].fillna(df[col].mode()[0])
```

```
In [ ]: print(df.columns)
```

```
Index(['id', 'age', 'sex', 'dataset', 'cp', 'trestbps', 'chol', 'fbs',
       'restecg', 'thalch', 'exang', 'oldpeak', 'slope', 'ca', 'thal',
       'num'],
      dtype='object')
```

```
In [ ]: # One-Hot Encoding for multi-class categorical features in UCI Heart Disease
dataset
```

```
multi_class_cols = ['cp', 'restecg', 'slope', 'ca', 'thal']

# Keep only valid columns
multi_class_cols = [col for col in multi_class_cols if col in df.columns]

if len(multi_class_cols) > 0:
    print("Applying One-Hot Encoding on:", multi_class_cols)

    from sklearn.preprocessing import OneHotEncoder
    import pandas as pd

    # One-Hot Encoder setup
    ohe = OneHotEncoder(sparse_output=False, drop='first') # updated for
    sklearn 1.2+ compatibility

    # Transform categorical columns
    ohe_array = ohe.fit_transform(df[multi_class_cols])

    # Convert to DataFrame
    ohe_df = pd.DataFrame(
        ohe_array,
        columns=ohe.get_feature_names_out(multi_class_cols),
        index=df.index
    )

    # Drop original categorical columns
    df = df.drop(multi_class_cols, axis=1)

    # Merge encoded features
    df = pd.concat([df, ohe_df], axis=1)

else:
    print("No multi-class categorical columns found for One-Hot Encoding.")
    df.info()
```

```
Applying One-Hot Encoding on: ['cp', 'restecg', 'slope', 'ca', 'thal']
```

```
In [ ]: from sklearn.preprocessing import StandardScaler

# Select numeric columns after encoding
numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns

# Remove ID column if exists
if 'id' in numeric_cols:
    numeric_cols = numeric_cols.drop('id')

# Remove target column
if 'num' in numeric_cols:
    numeric_cols = numeric_cols.drop('num')

print("Numeric Columns to Scale:", list(numeric_cols))

# Apply StandardScaler
scaler = StandardScaler()
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])

print("Scaling Completed!")
```

```
Numeric Columns to Scale: ['age', 'trestbps', 'chol', 'thalch',
 'oldpeak', 'cp_atypical angina', 'cp_non-anginal', 'cp_typical angina',
 'restecg_normal', 'restecg_st-t abnormality', 'restecg_nan',
 'slope_flat', 'slope_upsloping', 'slope_nan', 'ca_1.0', 'ca_2.0',
 'ca_3.0', 'ca_nan', 'thal_normal', 'thal_reversible defect',
 'thal_nan']
Scaling Completed!
```

```
In [ ]: from sklearn.model_selection import train_test_split
target = 'num'
# Remove ID and target from feature list
features = [c for c in df.columns if c not in ['id', 'num']]
# Split features and target
X = df[features]
y = df[target]
# Train-test split (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
print("Train shape:", X_train.shape, y_train.shape)
print("Test shape:", X_test.shape, y_test.shape)

print("\nTrain class distribution:\n", y_train.value_counts(normalize=True))
print("\nTest class distribution:\n", y_test.value_counts(normalize=True))
```

```
Train shape: (736, 25) (736,)
Test shape: (184, 25) (184,)
```

```
Train class distribution:
num
0    0.447011
1    0.288043
2    0.118207
3    0.116848
4    0.029891
Name: proportion, dtype: float64
```

```
Test class distribution:
num
0    0.445652
1    0.288043
2    0.119565
3    0.114130
4    0.032609
Name: proportion, dtype: float64
```

```
In [ ]: import pandas as pd
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
# BINARY CATEGORICAL COLUMNS → Label Encoding
binary_cols = ['sex', 'fbs', 'exang'] # update if needed
binary_cols = [c for c in binary_cols if c in df.columns]
if len(binary_cols) > 0:
    print("Label Encoding:", binary_cols)
    le = LabelEncoder()
    for col in binary_cols:
        df[col] = le.fit_transform(df[col].astype(str))
# MULTI-CLASS CATEGORICAL → One-Hot Encoding
multi_class_cols = ['cp', 'restecg', 'slope', 'ca', 'thal']
multi_class_cols = [col for col in multi_class_cols if col in df.columns]

if len(multi_class_cols) > 0:
    print("One-Hot Encoding:", multi_class_cols)

ohe = OneHotEncoder(sparse_output=False, drop='first')

ohe_array = ohe.fit_transform(df[multi_class_cols])

ohe_df = pd.DataFrame(
    ohe_array,
    columns=ohe.get_feature_names_out(multi_class_cols),
    index=df.index
)

df = df.drop(multi_class_cols, axis=1)
df = pd.concat([df, ohe_df], axis=1)
#Check final types
print("\nFinal columns type:")
print(df.dtypes)
```

Label Encoding: ['sex', 'fbs', 'exang']

	Final columns type:
id	int64
age	float64
sex	int64
dataset	object
trestbps	float64
chol	float64
fbs	int64
thalch	float64
exang	int64
oldpeak	float64
num	int64
cp_atypical angina	float64
cp_non-anginal	float64
cp_typical angina	float64
restecg_normal	float64
restecg_st-t abnormality	float64
restecg_nan	float64

```
slope_flat           float64
slope_upsloping      float64
slope_nan            float64
ca_1.0               float64
ca_2.0               float64
ca_3.0               float64
ca_nan               float64
thal_normal          float64
thal_reversible defect float64
thal_nan              float64
dtype: object
```

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier

df = pd.read_csv("/content/drive/MyDrive/project ai/UCI Heart Disease Data.csv")
df = df.loc[:, ~df.columns.str.contains('^\d+')]
df = df.fillna(df.median(numeric_only=True))
df['num'] = df['num'].apply(lambda x: 1 if x>0 else 0)

X = df.drop(['id','num'], axis=1)
X = pd.get_dummies(X, drop_first=True)    # <-
y = df['num']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

models = {
    "Logistic Regression": LogisticRegression(max_iter=500),
    "SVM": SVC(),
    "KNN": KNeighborsClassifier(),
    "Random Forest": RandomForestClassifier(),
    "Decision Tree": DecisionTreeClassifier()
}

for name, model in models.items():
    model.fit(X_train, y_train)
    pred = model.predict(X_test)
    acc = accuracy_score(y_test, pred)
    print(f"{name} Accuracy: {acc:.4f}")
```

```
Logistic Regression Accuracy: 0.8207
SVM Accuracy: 0.8478
KNN Accuracy: 0.8478
Random Forest Accuracy: 0.8424
Decision Tree Accuracy: 0.7880
```

```
In [ ]: from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
import seaborn as sns
import matplotlib.pyplot as plt

def evaluate_model(model, X_test, y_test, model_name):
    y_pred = model.predict(X_test)

    print(f"\n====")
    print(f"Model: {model_name}")
    print("====")

    print("Accuracy:", accuracy_score(y_test, y_pred))

    print("\nClassification Report:")
    print(classification_report(y_test, y_pred))


    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(5,4))
    sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
    plt.title(f"{model_name} - Confusion Matrix")
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

# RUN EVALUATION FOR ALL 5 MODELS
for name, model in models.items():
    evaluate_model(model, X_test, y_test, name)
```

```
=====
Model: Logistic Regression
=====
Accuracy: 0.8206521739130435

Classification Report:
precision    recall   f1-score   support
          0       0.84      0.74      0.79       82
          1       0.81      0.88      0.85      102

accuracy                           0.82      184
macro avg       0.82      0.81      0.82      184
weighted avg    0.82      0.82      0.82      184
```



```
=====
```


Model: SVM

```
=====
```

Accuracy: 0.8478260869565217

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.77	0.82	82
1	0.83	0.91	0.87	102
accuracy			0.85	184
macro avg	0.85	0.84	0.84	184
weighted avg	0.85	0.85	0.85	184



```
=====
```


Model: KNN

```
=====
```

Accuracy: 0.8478260869565217


Classification Report:

	precision	recall	f1-score	support
0	0.88	0.77	0.82	82
1	0.83	0.91	0.87	102
accuracy			0.85	184
macro avg	0.85	0.84	0.84	184
weighted avg	0.85	0.85	0.85	184




```
=====
Model: Random Forest
=====
Accuracy: 0.842391304347826
```

```
Classification Report:
precision    recall   f1-score   support
      0       0.84      0.80      0.82      82
      1       0.85      0.87      0.86     102
accuracy                           0.84      184
macro avg       0.84      0.84      0.84      184
weighted avg    0.84      0.84      0.84      184
```



```
=====
Model: Decision Tree
=====
Accuracy: 0.7880434782608695
```

```
Classification Report:
precision    recall   f1-score   support
      0       0.80      0.70      0.75      82
      1       0.78      0.86      0.82     102
accuracy                           0.79      184
macro avg       0.79      0.78      0.78      184
weighted avg    0.79      0.79      0.79      184
```





```
In [ ]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt



def plot_roc_curve(model, X_test, y_test, model_name):
    if hasattr(model, "predict_proba"):
        y_prob = model.predict_proba(X_test)[:, 1]
    elif hasattr(model, "decision_function"):
        y_prob = model.decision_function(X_test)
    else:
        print(f"{model_name} does not support ROC curve.")
        return


    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)

    plt.figure(figsize=(6, 5))
    plt.plot(fpr, tpr, label=f"{model_name} (AUC = {roc_auc:.3f})")
    plt.plot([0, 1], [0, 1], linestyle="--")
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.title(f"ROC Curve - {model_name}")
    plt.legend()
    plt.show()

# RUN ROC FOR ALL MODELS
for name, model in models.items():
    plot_roc_curve(model, X_test, y_test, name)
```








```
In [ ]: from sklearn.model_selection import learning_curve
import numpy as np
import matplotlib.pyplot as plt



def plot_learning_curve(model, X, y, model_name):
    train_sizes, train_scores, test_scores = learning_curve(
        model, X, y,
        cv=5,
        scoring='accuracy',
        train_sizes=np.linspace(0.1, 1.0, 10),
        n_jobs=-1
    )



    train_mean = np.mean(train_scores, axis=1)
    test_mean = np.mean(test_scores, axis=1)

    plt.figure(figsize=(6,5))
    plt.plot(train_sizes, train_mean, label="Training Accuracy")
    plt.plot(train_sizes, test_mean, label="Validation Accuracy")
    plt.title(f"Learning Curve - {model_name}")
    plt.xlabel("Training Set Size")
    plt.ylabel("Accuracy")
    plt.legend()
    plt.show()

# RUN LEARNING CURVES FOR ALL MODELS
for name, model in models.items():
    plot_learning_curve(model, X, y, name)
```







Exported with [runcell](#) — convert notebooks to HTML or PDF anytime at [runcell.dev](#).